

Understanding Machine Learning

Dr. Sunil Thomas T

College of Engineering Attingal

suniltt@gmail.com

<http://brainstorms.in>

<https://www.facebook.com/vu2swx>

August 20, 2016

Buzz words !!

- 1 Machine Learning
- 2 Data Mining
- 3 Pattern Recognition
- 4 Big Data
- 5 Recommender systems

Who uses it?

Lot of companies. In fact there is a growing requirement for ML/Data sciences professionals world wide.

Look at



Look at the ads. They are customized for you. The two web giants look at what you share or browse and learn about your habits, tastes and preferences. They then customize their advertisements based on this learning.

This talk ?

There are two parts in this workshop/

- I will give an informal introduction to machine learning in the first part
- Then I will introduce you to R, a popular tool used by machine learning/data sciences community. We will then write some simple machine learning programs in R.

I don't expect you to be hard core programmers. But a basic familiarity with programming is assumed. I will give some code that you can try right now.

Can Machines Learn ? Or can we teach Machines



How do we learn ?



How do we train kids ?



How do we train kids ?



How do we test our learning ?



How do we test our learning ?



What if we make errors ?



Retraining



Our Learning cycle

- Training
- Testing
- Retraining

How do we learn?

What kind of learning problems we encounter

- Classification
- Regression

How do we perform the above tasks.

- We learn using features of objects.

Feature ?

What are the features ?

Depends on the problem you are trying to solve.

- For example in the Bus vs Car problem we can have some of the following as features.
 - No of wheels
 - Size
 - No of windows
 -
 - Lot of other things

Now let us teach machines.

We need thousands of examples.

Features features and features

For each class of objects we need to find out good features

How do we represent features?

Table: Features for car

	Size	Weight	No of Wheels	No of Doors
Alto	8	700	4	4
i10	7.5	800	4	4
Honda City	9	950	4	4
Tata Nano	5	600	4	4


How do we represent features?

Table: Features for bus


	Size	Weight	No of Wheels	No of Doors
TATA	20	1700	6	2
Leyland	16	1800	6	2
Tata mini	23	1950	6	3
Volvo	28	1600	6	3

Feature Space


For simplicity I am showing a 3 dimensional feature space. But you can have as many dimensions as you want.



Objects in Feature Space




Classification Problem




The black line above represents a plane splitting the feature space into two.


Testing a classifier



Testing a classifier



Finding the best classifier-The Challenge



How complicated is the math ?

A bit if you are working on theoretical side.

If you are planning to ML algorithms, to your problems, it is very easy. There are lot of libraries, open source codes and sample data sets.



How complicated is the math ?

A bit if you are working on theoretical side.

If you are planning to ML algorithms, to your problems, it is very easy. There are lot of libraries, open source codes and sample data sets.




You need to know


- Linear Algebra
- Probability Theory
- Optimization Theory

Regression

You are give several data points. You have to find out a best fitting curve.



Types of Machine Learning



Sample data sets

First let us look at some real data that ML community uses. There are several sources.

Public data sets for machine learning.

UC Irvine Machine Learning Repository


Middlebury CV data set

Lot more.

iris data set



iris data set




iris data set

sepal		petal		class
length	width	length	width	
6.3	2.3	4.4	1.3	versicolor
6.2	3.4	5.4	2.3	virginica
5.2	3.4	1.4	0.2	setosa
6.9	3.1	5.4	2.1	virginica
5.7	4.4	1.5	0.4	setosa
5.4	3.7	1.5	0.2	setosa
5	3.3	1.4	0.2	setosa
6.4	2.8	5.6	2.1	virginica
6	3	4.8	1.8	virginica
5.5	2.5	4	1.3	versicolor

iris data set

3-D Scatterplot of Iris Data



Algorithms

- 1 k-Nearest Neighbors (kNN)
- 2 Support Vector Machines (SVM)
- 3 Expectation Maximization (EM)
- 4 Random forest
- 5 Naive Bayes
- 6 And many more

We will try out a few later

What are the opportunities

Research


Commercial

Let us look at some algorithms


k-Nearest Neighbor

Support vector machines


k-NN



Support vector machine



SVM kernel



Tools

- 1 weka
- 2 rapid miner
- 3 tensor flow
- 4 scikit learn
- 5 Many more

I guess tea and snacks are waiting

We can have a short break